

Computational Linguistics

CSC 2501 / 485
Fall 2018

9

Assignment Project Exam Help

9. Statistical parsing

<https://powcoder.com>

Add WeChat powcoder

Gerald Penn

Department of Computer Science, University of Toronto

Reading: Jurafsky & Martin: 5.2–5.5.2, 5.6, 12.4, 14.0–1, 14.3–4, 14.6–7. Bird et al: 8.6.

Copyright © 2017 Suzanne
Stevenson, Graeme Hirst
and Gerald Penn. All rights
reserved.

Statistical parsing 1

- **General idea:**
 - Assign probabilities to rules in a context-free grammar.
 - Use a likelihood model. [Assignment Project Exam Help](#)
 - Combine probabilities of rules in a tree. <https://powcoder.com>
 - Yields likelihood of a parse. [Add WeChat powcoder](#)
 - The best parse is the *most likely* one.

Statistical parsing 2

- ***Motivations:***

- Uniform process for attachment decisions.
- Use lexical preferences in all decisions.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Two general approaches

1. Assign a probability to each rule of grammar, including lexical productions.
 - Parse string of input words with probabilistic rules.
The can will rust.
2. Assign probabilities only to non-lexical productions.
 - Probabilistically tag input words with syntactic categories using a **part-of-speech tagger**.
 - Consider the pre-terminal syntactic categories to be terminals, parse that string with probabilistic rules.
Det N Modal Verb.
3. “Supertagging” – parsing as tagging with tree fragments.

Part-of-speech tagging 1

- ***Part-of-speech (PoS) tagging:***
Given a sequence of words $w_1 \dots w_n$ (from well-formed text), determine the syntactic category (PoS) C_i of each word.
- */i.e., the best category sequence $C_1 \dots C_n$ to assign to the word sequence $w_1 \dots w_n$.*

Most likely

Part-of-speech tagging 2

- Example:

The can will rust
det modal verb modal verb noun
noun noun verb
verb verb

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Part-of-speech tagging 3

$$P(C_1 \dots C_n | w_1 \dots w_n) = \frac{P(C_1 \dots C_n \wedge w_1 \dots w_n)}{P(w_1 \dots w_n)}$$

- We cannot get this probability directly.
- Have to estimate it (through counts).
- Perhaps after first approximating it (by modifying the formula).
- Counts: Need representative corpus.

PoS tagging: Unigram MLE 1

- Look at individual words (***unigrams***):

$$P(C|w) = \frac{P(C \wedge w)}{P(w)}$$

- Maximum likelihood estimator (MLE):

<https://powcoder.com>

$$P(C|w) = \frac{c(w \text{ is } C)}{c(w)}$$

Count in corpus

PoS tagging: Unigram MLE 2

- Problems of MLE:
 - Sparse data.
 - Extreme cases:
 - a. Undefined if w is not in the corpus.
 - b. 0 if w does not appear in a particular category.
- <https://powcoder.com>
- Add WeChat powcoder

PoS tagging: Unigram MLE 3

- **Smoothing** of formula, e.g.,:
$$P(C|w) \approx \frac{c(w \text{ is } C) + \epsilon}{c(w) + \epsilon N}$$
- Give small (non-zero) probability value to unseen events, taken from seen events by *discounting* them. <https://powcoder.com>
Assignment Project Exam Help
Add WeChat powcoder
- Various methods to ensure we still have valid probability distribution.

PoS tagging: Unigram MLE 4

- Just choosing the most frequent PoS for each word yields 90% accuracy in PoS tagging.
- But:
 - Not uniform across words.
 - Accuracy is low (0.9^n) when multiplied over n words.
 - No context: *The fly* vs. *I will fly*.
- Need better approximations for

$$P(C_1 \dots C_n | w_1 \dots w_n)$$

PoS tagging: Bayesian method

- Use Bayes's rule to rewrite:

$$P(C_1 \dots C_n | w_1 \dots w_n) = \frac{P(C_1 \dots C_n) \times P(w_1 \dots w_n | C_1 \dots C_n)}{P(w_1 \dots w_n)}$$

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

- For a given word string, we want to maximize this, find most likely $C_1 \dots C_n$:

$$\operatorname{argmax}_{C_1 \dots C_n} P(C_1 \dots C_n | w_1 \dots w_n)$$

- So just need to maximize the numerator.

Approximating probabilities 1

- Approximate $P(C_1 \dots C_n)$ by predicting each category from previous $n-1$ categories: an ***n*-gram model**.

Warning: Not the same *n*!!

- Bigram (2-gram) model:

$$P(C_1 \dots C_n) \approx \prod_{i=1}^n P(C_i | C_{i-1})$$

<https://powcoder.com>
Add WeChat powcoder

- Posit pseudo-categories START at C_0 , and END as C_n . Example:

$$P(A \ N \ V \ N) \approx P(A|\text{START}) \cdot P(N|A) \cdot P(V|N) \cdot P(\text{END}|N)$$

Approximating probabilities 2

- Approximate ② $P(w_1 \dots w_n | C_1 \dots C_n)$ by assuming that the probability of a word appearing in a category is independent of the words surrounding it.

Assignment Project Exam Help

$$P(w_1 \dots w_n | C_1 \dots C_n) \approx \prod_{i=1}^n P(w_i | C_i)$$

<https://powcoder.com>

Add WeChat powcoder

Lexical generation
probabilities

Approximating probabilities 3

- Why is $P(w|C)$ better than $P(C|w)$?
 - $P(C|w)$ is clearly *not* independent of surrounding categories.
 - Lexical generation probability is somewhat more independent. <https://powcoder.com>
 - Complete formula for PoS includes bigrams, and so it does capture some context.

Putting it all together

$$\begin{aligned}
 & P(C_1 \dots C_n \mid w_1 \dots w_n) \\
 &= \frac{P(C_1 \dots C_n \wedge w_1 \dots w_n)}{P(w_1 \dots w_n)} \\
 &= \frac{P(C_1 \dots C_n) \times P(w_1 \dots w_n \mid C_1 \dots C_n)}{P(w_1 \dots w_n)} \\
 &\propto P(C_1 \dots C_n) \times P(w_1 \dots w_n \mid C_1 \dots C_n) \\
 &\approx \prod_{i=1}^n P(C_i \mid C_{i-1}) \times P(w_i \mid C_i) \quad \text{③} \\
 &= \left[\prod_{i=1}^n \frac{c(C_{i-1}C_i)}{c(C_{i-1})} \times \frac{c(w_i \text{ is } C_i)}{c(C_i)} \right]
 \end{aligned}$$

Really should use smoothed MLE; MLE for categories not the same as for words;
 cf slide 10 cf slide 8

Finding max 1

- Want to find the argmax (most probable) $C_1 \dots C_n$.
- *Brute force method*: Find all possible sequences of categories and compute P .
[Assignment Project Exam Help](https://powcoder.com)
- Unnecessary: <https://powcoder.com> Our approximation assumes independence: [Add WeChat powcoder](#)
 - *Category bigrams*: C_i depends only on C_{i-1} .
 - *Lexical generation*: w_i depends only on C_i .
- Hence we do not need to enumerate all sequences independently.

Finding max 2

- Bigrams:
Markov model.
- States are categories and transitions are bigrams.
- Lexical generation probabilities:
Hidden Markov model.
- Words are outputs (with given probability) of states.
- A word could be the output of more than one state.
- Current state is unknown (“hidden”).

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Example

Based on an example in section 7.3 of: Allen, James. *Natural Language Understanding* (2nd ed), 1995, Benjamin Cummings.

- Artificial corpus of PoS-tagged 300 sentences using only Det, N, V, P.

- *The flower flowers like a bird.*
Some birds like a flower with fruit beetles.
Like flies like flies.

...

<https://powcoder.com>


- Some lexical generation probabilities:

$P(\text{the} \text{Det}) = .54$	$P(\text{like} \text{N}) = .012$	$P(\text{flower} \text{N}) = .063$	$P(\text{birds} \text{N}) = .076$
$P(\text{a} \text{Det}) = .36$	$P(\text{like} \text{V}) = .1$	$P(\text{flower} \text{V}) = .050$	$P(\text{flies} \text{V}) = .076$
$P(\text{a} \text{N}) = .001$	$P(\text{like} \text{P}) = .068$	$P(\text{flowers} \text{N}) = .050$	$P(\text{flies} \text{N}) = .025$
\vdots	\vdots	$P(\text{flowers} \text{V}) = .053$	\vdots
		\vdots	

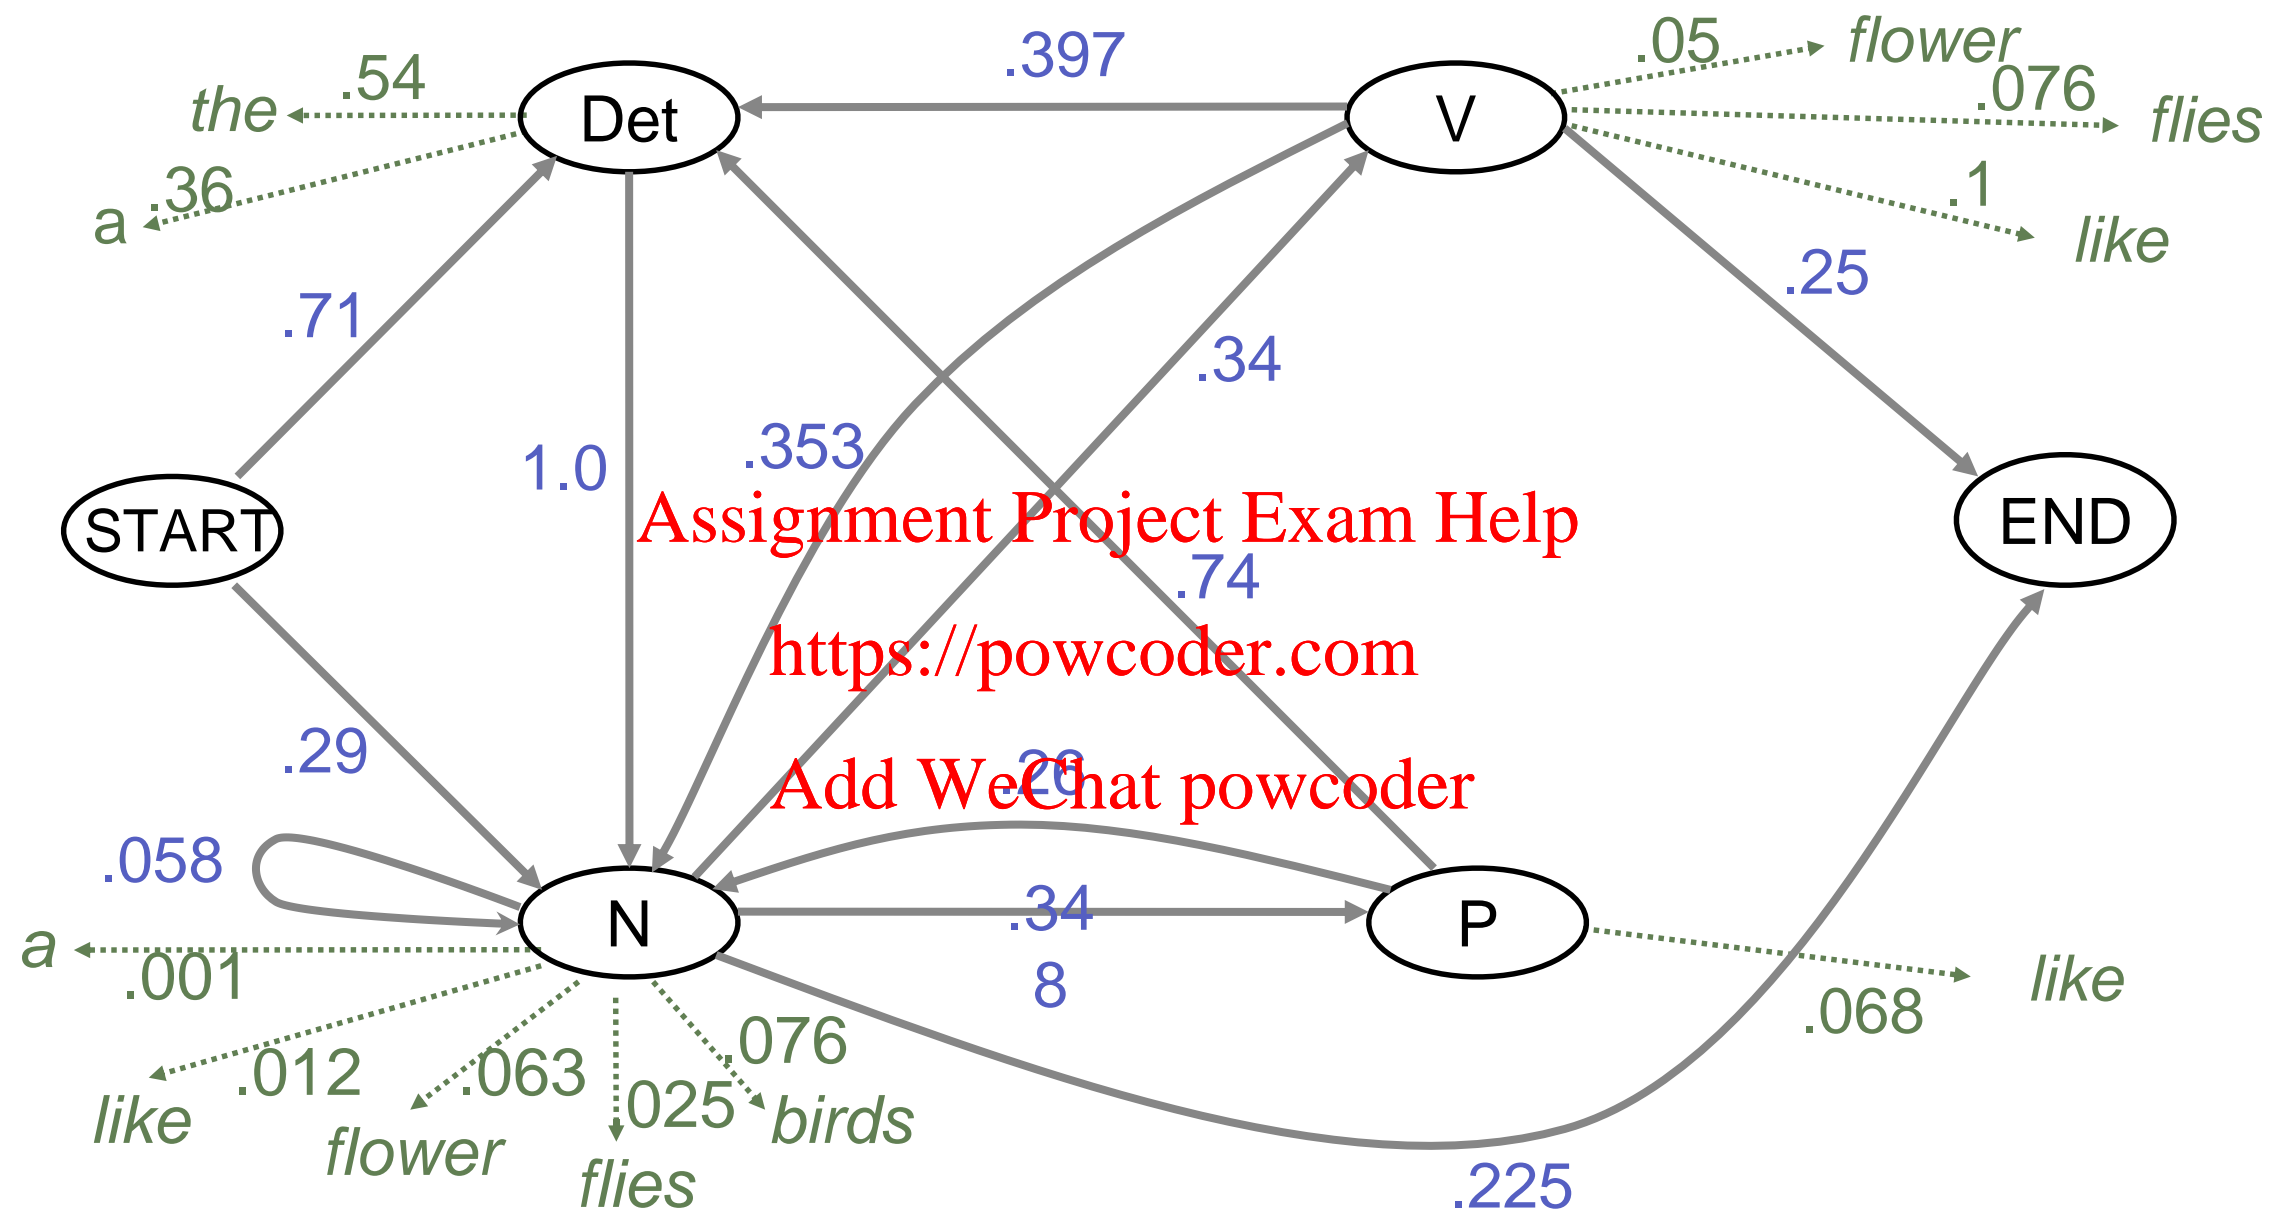
Markov model: Bigram table

Bigram C_{i-1}, C_i	Count C_{i-1}	Count C_{i-1}, C_i	$P(C_i C_{i-1})$	Estimate
START, Det	300	213	$P(\text{Det} \text{START})$	0.710
START, N	300	87	$P(\text{N} \text{START})$	0.290
Det, N	558	558	$P(\text{N} \text{Det})$	1.000
N, V	883	300	$P(\text{V} \text{N})$	0.340
N, N	883	51	$P(\text{N} \text{N})$	0.058
N, P	883	307	$P(\text{P} \text{N})$	0.348
N, END	883	225	$P(\text{END} \text{N})$	0.255
V, N	300	106	$P(\text{N} \text{V})$	0.353
V, Det	300	119	$P(\text{Det} \text{V})$	0.397
V, END	300	75	$P(\text{END} \text{V})$	0.250
P, Det	307	226	$P(\text{Det} \text{P})$	0.740
P, N	307	81	$P(\text{N} \text{P})$	0.260

Markov model: Transition probabilities



HMM: Lexical generation probabilities



$P(the Det) = .54$	$P(like N) = .012$	$P(flower N) = .063$	$P(birds N) = .076$
$P(a Det) = .36$	$P(like V) = .1$	$P(flower V) = .050$	$P(flies V) = .076$
$P(a N) = .001$	$P(like P) = .068$	$P(flowers N) = .050$	$P(flies N) = .025$
\vdots	\vdots	$P(flowers V) = .053$	\vdots

Hidden Markov models 1

- Given the observed output, we want to find the most likely path through the model.

The can will rust

det modal verb **modal verb** noun
noun noun verb
verb verb

Hidden Markov models 2

- At any state in an HMM, how you got there is irrelevant to computing the next transition.
- So, just need to remember the best path and probability up to that point.
Assignment Project Exam Help
- Define $P_{C_{i-1}}$ as the probability of the best sequence up to state C_{i-1} .
<https://powcoder.com>
Add WeChat powcoder
- Then find C_i that maximizes
 $P(C_{i-1}) \times P(C_i|C_{i-1}) \times P(w|C_i)$ $P(C_i|C_{i-1}) \times P(w|C_i)$ 3 from slide 17

Viterbi algorithm

- Given an HMM and an observation O of its output, finds the most probable sequence S of states that produced O .
- O = words of sentence, S = PoS tags of sentence
- Parameters of HMM based on large training corpus.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Statistical chart parsing 1

- Consider tags as terminals (*i.e.*, use a PoS tagger to pre-process input texts).
Det N Modal Verb.
- For probability of each grammar rule, use MLE.
<https://powcoder.com>
- Probabilities derived from hand-parsed corpora (treebanks).
[Add WeChat powcoder](#)
- Count frequency of use c of each rule $C \rightarrow \alpha$, for each non-terminal C and each different RHS α .

What are some problems with this approach?

Statistical chart parsing 2

- MLE probability of rules:

- For each rule $C \rightarrow \alpha$:

$$P(C \rightarrow \alpha | C) = \frac{c(C \rightarrow \alpha)}{\sum_{\beta} c(C \rightarrow \beta)} = \frac{c(C \rightarrow \alpha)}{c(C)} \quad 4$$

- Takes no account of the context of use of a rule: ***independence assumption***.
- Source-normalized***: assumes a top-down generative process.
- NLTK's pchart demo doesn't POS-tag first (words are generated top-down), and it shows $P(t)$ rather than $P(t|s)$. Why?

```
>>> import nltk
>>> nltk.parse.pchart.demo()
```

```
1: I saw John with my telescope
   <Grammar with 17 productions>
```

```
2: the boy saw Jack with Bob under the table with a telescope
   <Grammar with 23 productions>
```

Which demo (1-2)? 1

```
s: I saw John with my telescope
```

```
parser: <nltk.parse.pchart.InsideChartParser object at 0x7f61288f3290>
```

```
grammar: Grammar with 17 productions (start state = S)
```

```
S -> NP VP [1.0]
```

```
NP -> Det N [0.5]
```

```
NP -> NP PP [0.25]
```

```
NP -> 'John' [0.1]
```

```
NP -> 'I' [0.15]
```

```
Det -> 'the' [0.8]
```

```
Det -> 'my' [0.2]
```

```
N -> 'man' [0.5]
```

```
N -> 'telescope' [0.5]
```

```
VP -> VP PP [0.1]
```

```
VP -> V NP [0.7]
```

```
VP -> V [0.2]
```

```
V -> 'ate' [0.35]
```

```
V -> 'saw' [0.65]
```

```
PP -> P NP [1.0]
```

```
P -> 'with' [0.61]
```

```
P -> 'under' [0.39]
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

[-]	[0:1]	'I'	[1.0]
. [-]	[1:2]	'saw'	[1.0]
. . [-]	[2:3]	'John'	[1.0]
. . . [-]	[3:4]	'with'	[1.0]
. . . . [-]	[4:5]	'my'	[1.0]
. [-]	[5:6]	'telescope'	[1.0]
. [-]	[5:6]	'telescope'	[1.0]
. . . . [-]	[4:5]	'my'	[1.0]
. . . [-]	[3:4]	'with'	[1.0]
. . [-]	[2:3]	'John'	[1.0]
. [-]	[1:2]	'saw'	[1.0]
[-]	[0:1]	'I'	[1.0]
. [-]	[1:2]	V -> 'saw' *	[0.65]
. >	[1:1]	VP -> * V NP	[0.7]
. >	[1:1]	V -> 'saw'	[0.65]
. . . [-]	[3:4]	P -> 'with' *	[0.61]
. . . >	[3:3]	PP -> * P NP	[1.0]
. . . [->	[3:4]	PP -> P * NP	[0.61]
. . . >	[3:3]	P -> * 'with'	[0.61]
. [-]	[5:6]	N -> 'telescope' *	[0.5]
. >	[5:5]	N -> * 'telescope'	[0.5]
. [->	[1:2]	VP -> V * NP	[0.455]
. >	[1:1]	VP -> * V	[0.2]
. . . . [-]	[4:5]	Det -> 'my' *	[0.2]
. . . . >	[4:4]	NP -> * Det N	[0.5]
. . . . >	[4:4]	Det -> * 'my'	[0.2]

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

⋮

```

      :
      :
| . . . . > . . | [4:4] S -> * NP VP [1.0]
| . . . . > . . | [4:4] NP -> * NP PP [0.25]
| . . . . [---> | [4:6] S -> NP * VP [0.05]
| . [---] . . . | [1:3] VP -> V NP * [0.0455]
| [-> . . . . . | [0:1] NP -> NP * PP [0.0375]
| . . . [-----] | [3:6] PP -> P NP * [0.0305]
| . . [-> . . . | [2:3] NP -> NP * PP [0.025]
| [---] . . . . | [0:2] S -> NP VP * [0.0195]
| . [-> . . . . | [1:2] VP -> VP * PP [0.013]
| . . . . [---> | [4:6] NP -> NP * PP [0.0125]
| [-----] . . . | [0:3] S -> NP VP * [0.006825]
| . [---> . . . | [1:3] VP -> VP * PP [0.00455]
| . . [-----] | [2:6] NP -> NP PP * [0.0007625]
| . . [-----> | [2:6] S -> NP * VP [0.0007625]
| . [-----] | [1:6] VP -> V NP * [0.0003469375]
| . . [-----> | [2:6] NP -> NP * PP [0.000190625]
| . [-----] | [1:6] VP -> VP PP * [0.000138775]
| [=====] | [0:6] S -> NP VP * [5.2040625e-05]
| . [-----> | [1:6] VP -> VP * PP [3.469375e-05]
| [=====] | [0:6] S -> NP VP * [2.081625e-05]
| . [-----> | [1:6] VP -> VP * PP [1.38775e-05]

```

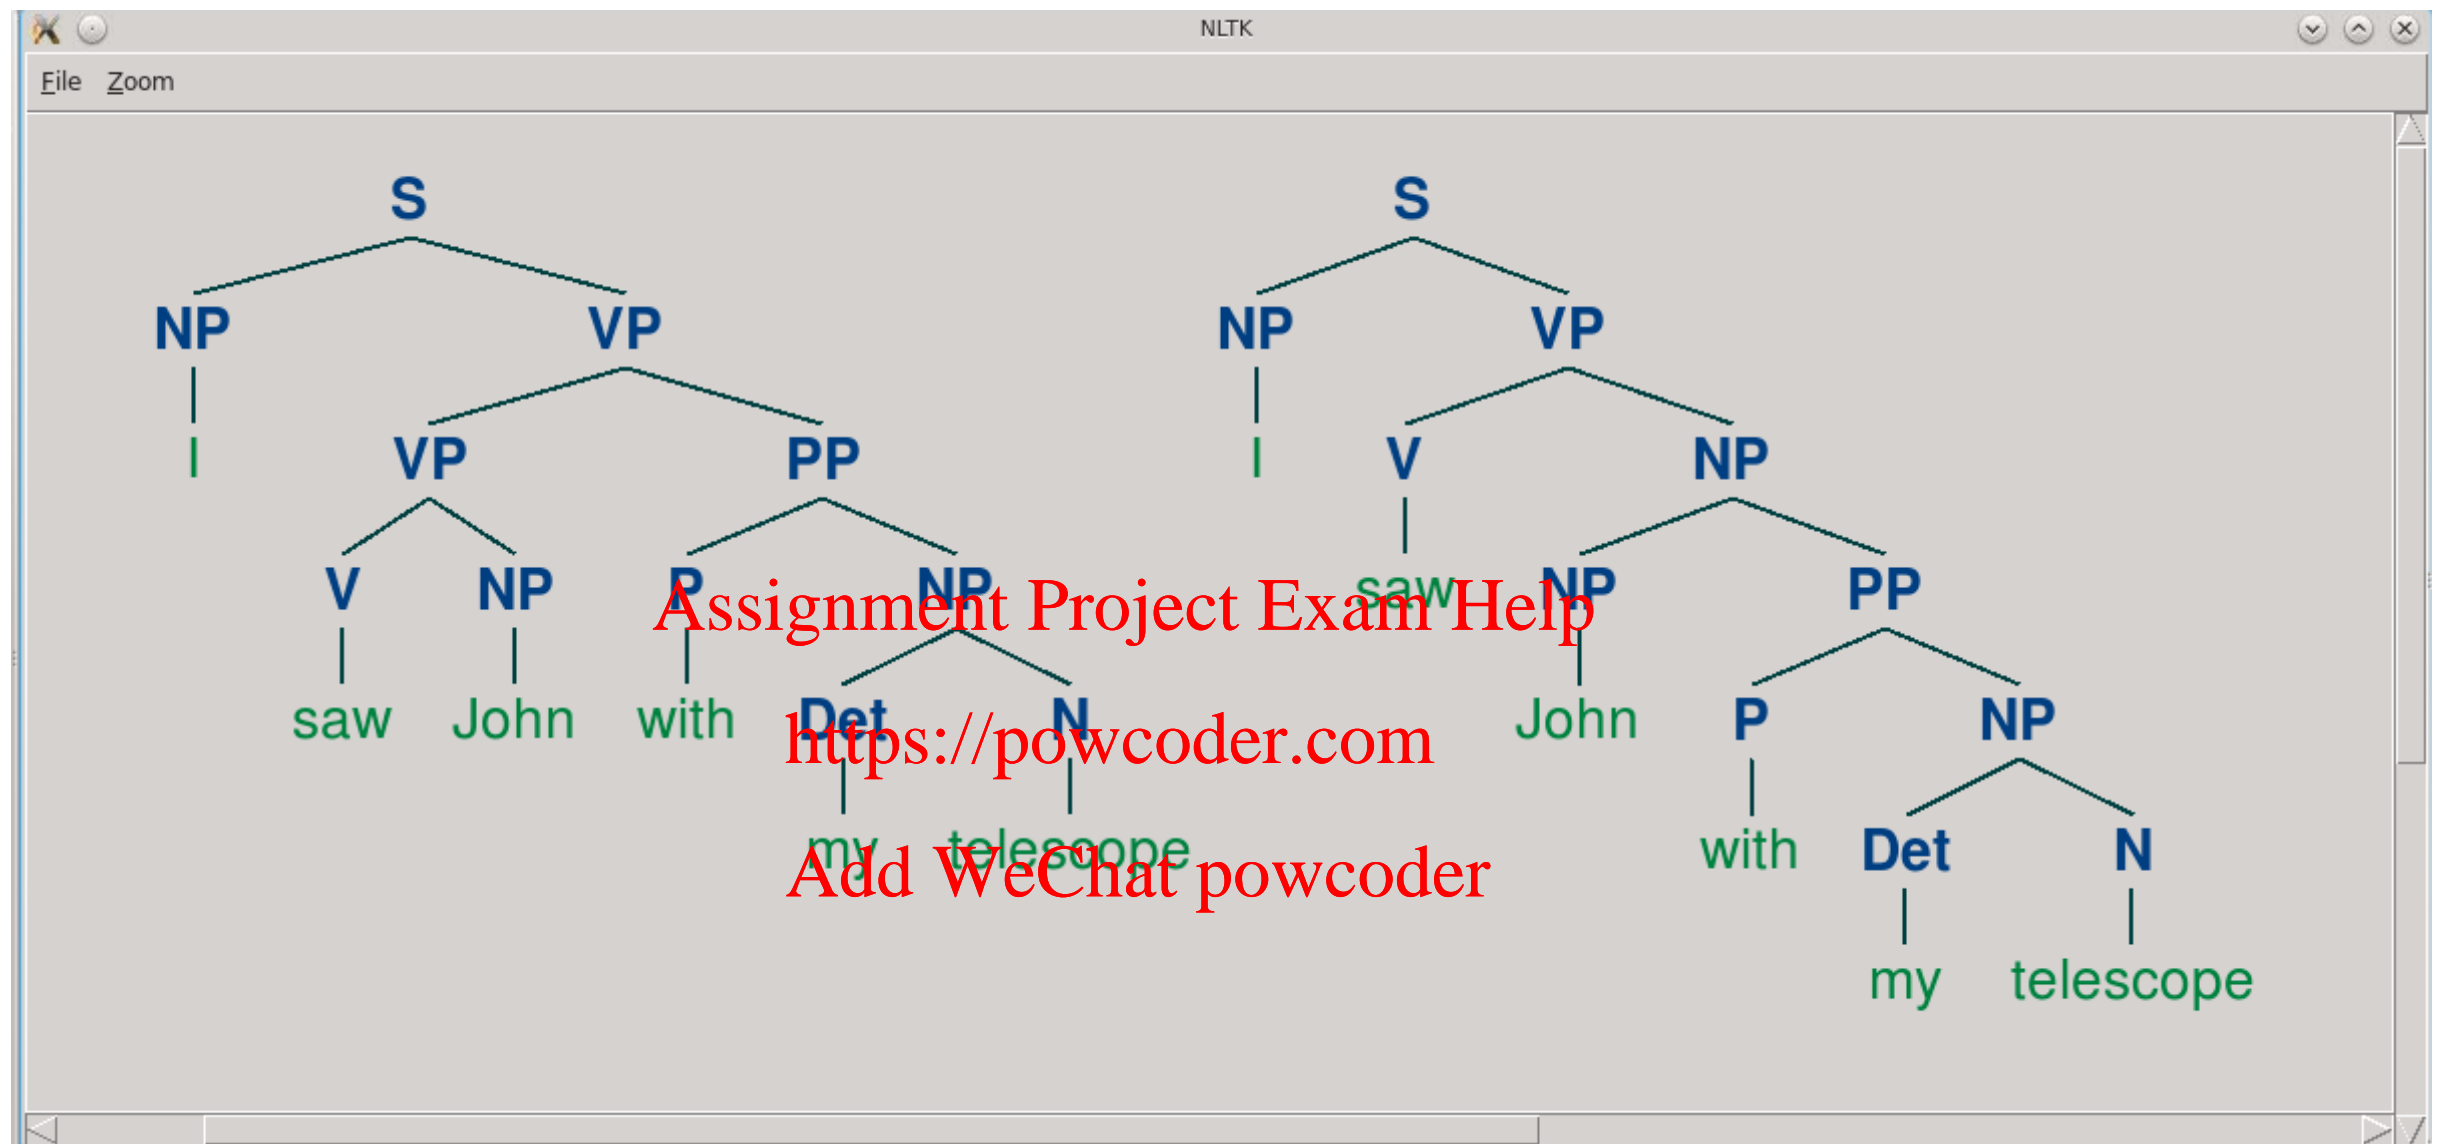
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Draw parses (y/n)? y
please wait...



Print parses (y/n)? y

```
(S
  (NP I)
  (VP
    (VP (V saw) (NP John))
    (PP (P with) (NP (Det my) (N telescope))))) [2.081625e-05]
```

```
(S
  (NP I)
  (VP
    (V saw)
    (NP
      (NP John)
      (PP (P with) (NP (Det my) (N telescope))))) [5.2040625e-05]
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Statistical chart parsing 3

- In this view of chart parsing, probability of chart entries is relatively simple to calculate. For completed constituents:

$$\begin{aligned} P(e_0) &= P(C_0 \rightarrow C_1 \dots C_n | C_0) \times P(e_1) \times \dots \times P(e_n) \\ &= P(C_0 \rightarrow C_1 \dots C_n | C_0) \times \prod_{i=1}^n P(e_i) \end{aligned}$$

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

5

e_0 is the entry for current constituent, of category C_0 ;
 $e_1 \dots e_n$ are chart entries for $C_1 \dots C_n$ in the RHS of the rule.

NB: Unlike for PoS tagging above, the C_i are not necessarily lexical categories.

Statistical chart parsing 4

- Consider a complete parse tree, t , with root label S .

- Recasting 5, t has the probability:

$$P(t) = P(S) \cdot \prod_n P(\text{rule}(n) | \text{cat}(n))$$

where n ranges over all nodes in the tree t ,

$\text{rule}(n)$ is the rule used for n ;

$\text{cat}(n)$ is the category of n .

- $P(S) = 1!$
- “Bottoms out” at lexical categories.
- Note that we’re parsing bottom-up, but the generative model “thinks” top-down regardless.

6

Inside-Outside Algorithm

- Maximum likelihood estimates on an annotated corpus can be improved to increase the likelihood of a different, unannotated corpus
- Step 1: parse the unannotated corpus using the MLE parameters
- Step 2: adjust the parameters according to the expected relative frequencies of different rules in the parse trees obtained in Step 1:
 - $\hat{p}(A \rightarrow B \ C) = \mu(A \rightarrow B \ C) / Z$
 - $\hat{p}(A \rightarrow w) = \mu(A \rightarrow w) / Z$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Inside-Outside Algorithm 2

- $\mu(A \rightarrow BC) = \sum_{\{i,k,j\}} \mu(A \rightarrow BC, i, k, j)$
- $\mu(A \rightarrow w) = \sum_i \mu(A, i) \delta_i(w)$

where we now count having seen an A from i to j, a B from i to k, and a C from k to j,
...or an A at location i, where there appears the word w).

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Inside-Outside Algorithm 3

- We can define these position-specific μ 's in terms of:

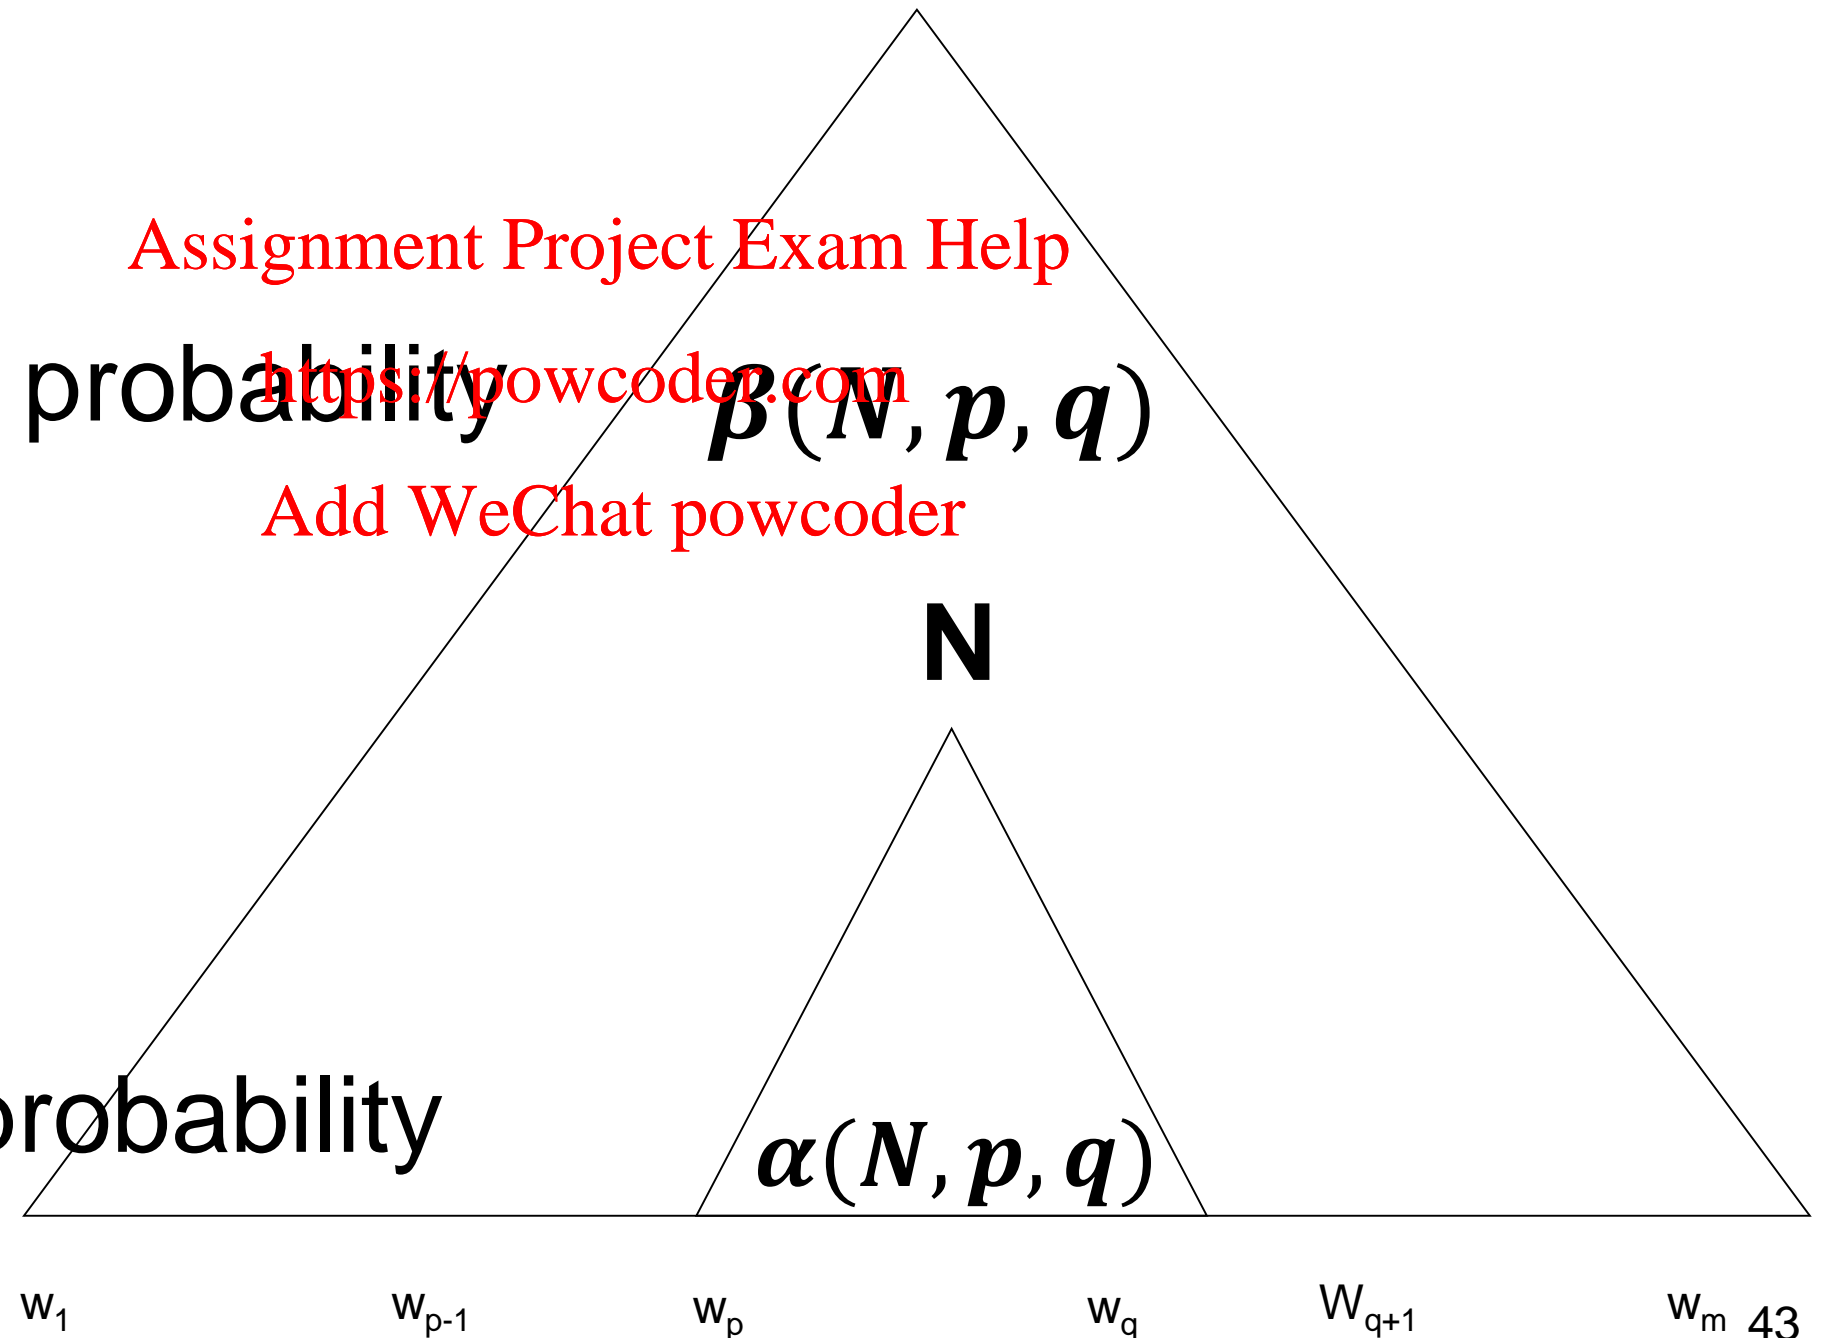
- outside probability $\beta(N, p, q)$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- inside probability



Inside-Outside Algorithm 4

- $\mu(A \rightarrow BC, i, k, j) = p(A \rightarrow BC) \beta(A, i, j) \alpha(B, i, k) \alpha(C, k + 1, j)$
 - $\mu(A, i) = \mu(A, i, i)$
 - $\mu(A, i, j) = \alpha(A, i, j) \beta(A, i, j)$
 - $Z = \alpha(S, 1, n)$
- Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

There are also very terse, recursive formulations of α and β that are amenable to dynamic programming.

Statistical chart parsing 5

- But just like non-statistical chart parsers, this one only answers ‘yes’ or ‘no’ (with a probability) in polynomial time:
 - It's not supposed to matter how we got each constituent. Just the non-terminal label and the span are all that should matter.
<https://powcoder.com>
Add WeChat powcoder
- There might be exponentially many trees in this formulation.
- And we're not calculating the probability that the input is a sentence – this is only the probability of one interpretation (tree).

Evaluation 1

- Evaluation method:
 - *Train* on part of a parsed corpus.
(*I.e.*, gather rules and statistics.)
 - Test on a different part of the corpus.
- In one sense, the best evaluation of a method like this would be data likelihood, but since we're scoring trees instead of strings, it's difficult to defend any sort of intuition about the numbers assigned to them.

Evaluation 2

- **Evaluation:** PARSEVAL measures compare parser output to known correct parse:
 - Labelled **precision**, labelled **recall**.
 - Assignment Project Exam Help
<https://powcoder.com>
Fraction of constituents in output that are correct.
 - Add WeChat powcoder
Fraction of correct constituents in output.
 - **F-measure** = harmonic mean of precision and recall = $2PR / (P + R)$

Evaluation 3

- **Evaluation:** PARSEVAL measures compare parser output to known correct parse:
- Penalize for **cross-brackets** per sentence:
Constituents in output that overlap two (or more) correct ones; e.g., `[[A B] C]` for `[A [B C]]`.
Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder
`[[Nadia] [smelled] [the eggplant]]`
`[[[Nadia] [smelled]] [the eggplant]]`

The labels on the subtrees aren't necessary for this one.

Evaluation 4

- PARSEVAL is a *classifier accuracy* score – much more extensional. All that matters is the right answer at the end.
- But that still means that we can look at parts of the right answer <https://powcoder.com>
- Can get ~75% labelled precision, recall, and F with above methods. [Add WeChat powcoder](#)

Improving statistical parsing

- Problem: Probabilities are based only on structures and categories:

$$P(C \rightarrow \alpha | C) = \frac{c(C \rightarrow \alpha)}{\sum_{\beta} c(C \rightarrow \beta)} = \frac{c(C \rightarrow \alpha)}{c(C)} \quad 4$$

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

- But actual words strongly condition which rule is used (*cf* Ratnaparkhi).
- Improve results by conditioning on more factors, including words. Think *semantics* – the words themselves give us a little bit of access to this.

Lexicalized grammars 1

- **Head** of a phrase: its central or key word.
 - The noun of an NP, the preposition of a PP, etc.
- **Lexicalized** grammar: Refine the grammar so that rules take heads of phrases into account — the actual words.
 - BEFORE: Rule for NP.
AFTER: Rules for NP-whose-head-is-*aardvark*, NP-whose-head-is-*abacus*, ..., NP-whose-head-is-*zymurgy*.
- And similarly for VP, PP, etc.

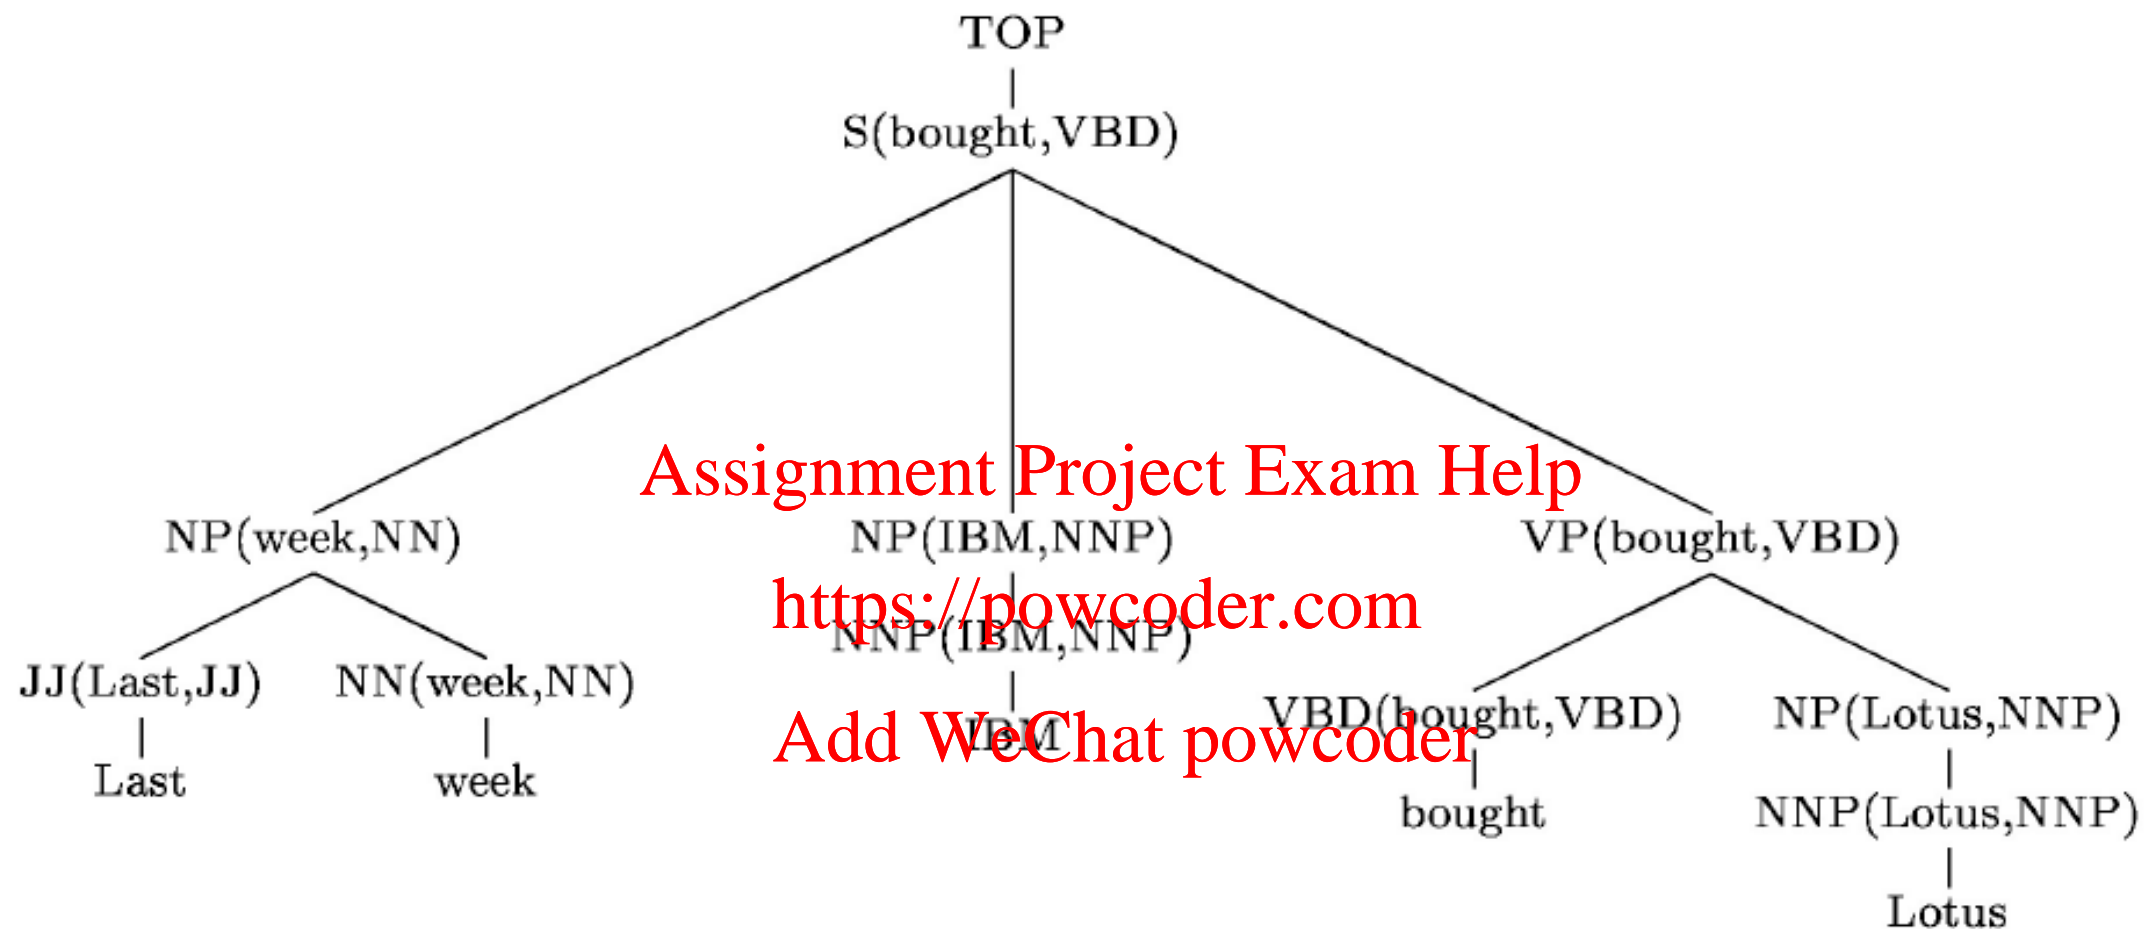
Lexicalized grammars 2

- Notation: $cat(head, tag)$ for constituent category cat headed by $head$ with part-of-speech tag .
- e.g., NP(*aardvark*, NN), PP(*without*, IN)

NP-whose-head-is-the-NN-*aardvark*

PP-whose-head-is-the-IN-*without*

A lexicalized grammar



TOP \rightarrow S(*bought*,VBD)

S(*bought*,VBD) \rightarrow NP(*week*,NN) NP(*IBM*,NNP)

VP(*bought*,VBD)

NP(*week*,NN) \rightarrow JJ(*Last*,JJ) NN(*week*,NN)

NP(*IBM*,NNP) \rightarrow NNP(*IBM*,NNP)

VP(*bought*,VBD) \rightarrow VBD(*bought*,VBD)

NP(*Lotus*,NNP)

NP(*Lotus*,NNP) \rightarrow NNP(*Lotus*,NNP)

Lexical Rules:

JJ(*Last*,JJ) \rightarrow *Last*

NN(*week*,NN) \rightarrow *week*

NNP(*IBM*,NNP) \rightarrow *IBM*

VBD(*bought*,VBD) \rightarrow *bought*

NNP(*Lotus*,NNP) \rightarrow *Lotus*

Lexicalized grammars 3

- Number of rules and categories explodes, but no theoretical change in parsing process (whether statistical or not).
- But far too specific for practical use; each is too rarely used to determine its probability.
<https://powcoder.com>
- Need something more than regular (unlexicalized) rules and less than complete lexicalization ...
[Add WeChat powcoder](#)
- ... perhaps we should change the process after all.

Lexicalized parsing 1

Starting from unlexicalized rules:

- **1. Lexicalization:** Consider the head word of each node, not just its category:

- $$P(t) = P(S) * \prod_n P(\text{rule}(n) | \text{head}(n))$$

Replaces 6
from slide 40

where $\text{head}(n)$ is the PoS-tagged head word of node n .

- Needs finer-grained probabilities:
 - e.g., probability that rule r is used, given we have an NP whose head is the noun *deficit*.

Lexicalized parsing 2

- **2. Head and parent:** Condition on the head and the head of the parent node in the tree:

$$P(\text{Sentence}, \text{Tree}) = \prod_{n \in \text{Tree}} P(\text{rule}(n) \mid \text{head}(n)) \times P(\text{head}(n) \mid \text{head}(\text{parent}(n)))$$

e.g., probability of rule r given that head is the noun *deficit*.

e.g., probability that head is the noun *deficit*,
given that parent phrase's head is the verb *report*.

Effects on parsing

- Lexical information introduces *context* into CFG.
- Grammar is larger.
- Potential problems of sparse data.
 - Possible solutions: Smoothing; back-off estimates.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

If you don't have data for a fine-grained situation, use data from a coarser-grained situation that it's contained in.

Bikel's 2004 interpretation

- Can condition on *any* information available in generating the tree.
- **Basic idea:** Avoid sparseness of lexicalization by decomposing rules.
[Assignment Project Exam Help](https://powcoder.com)
- Make plausible independence assumptions.
<https://powcoder.com>
- Break rules down into small steps (small number of parameters).
[Add WeChat powcoder](https://powcoder.com)
- Each rule still parameterized with word/PoS pair:
 $S(bought, VBD) \rightarrow NP(week, NN) NP(IBM, NNP) VP(bought, VBD)$

Collins's "model 1" 1

- **Lexical Rules**, with probability 1:
 $tag(word, tag) \rightarrow word$
- **Internal Rules**, with treebank-based probabilities. Separate terminals to the left and right of the head; generate one at a time:

Add WeChat powcoder
<https://powcoder.com>

$$X \rightarrow L_n L_{n-1} \dots L_1 H R_1 \dots R_{m-1} R_m \quad (n, m \geq 0)$$

X , L_i , H , and R_i all have the form $cat(head, tag)$.
Notation: Italic lowercase symbol for $(head, tag)$:

$$X(x) \rightarrow L_n(l_n) L_{n-1}(l_{n-1}) \dots L_1(l_1) H(h) R_1(r_1) \dots R_{m-1}(r_{m-1}) R_m(r_m)$$

Collins's “model 1” 2

- Assume there are additional L_{n+1} and R_{m+1} representing phrase boundaries (“STOP”).
- Example:
 $S(bought, VBD) \rightarrow NP(\text{week}, NN) NP(IBM, NNP) VP(bought, VBD)$
 $n = 2, m = 0$ (two constituents on the left of the head, zero on the right).
 $X = S, H = VP, L_1 = NP, L_2 = NP, L_3 = STOP, R_1 = STOP.$
 $h = (bought, VBD), l_1 = (IBM, NNP), l_2 = (week, NN).$
- Distinguish probabilities of heads P_h , of left constituents P_l , and of right constituents P_r .

Probabilities of internal rules 1

$$\begin{aligned} P(X(h)) &= P(L_{n+1}(l_{n+1}) L_n(l_n) \dots L_1(l_1) H(h) R_1(r_1) \dots R_m(r_m) R_{m+1}(r_{m+1}) | X, h) \\ &= P_h(H | X, h) \\ &\quad \times \prod_{i=1}^{n+1} P_l(L_i(l_i) | L_1(l_1) \dots L_{i-1}(l_{i-1}), X, h, H) \\ &\quad \times \prod_{j=1}^{m+1} P_r(R_j(r_j) | L_1(l_1) \dots L_n(l_n), R_1(r_1) \dots R_{j-1}(r_{j-1}), X, h, H) \end{aligned}$$

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

$$\approx P_h(H | X, h) \times \prod_{i=1}^{n+1} P_l(L_i(l_i) | X, h, H) \times \prod_{j=1}^{m+1} P_r(R_j(r_j) | X, h, H)$$

Generate head constituent

Generate left modifiers (stop at STOP)

By independence assumption

Generate right modifiers (stop at STOP)

Probabilities of internal rules 2

Example:

$P(S(bought, VBD))$
 $\rightarrow NP(\text{week}, NN) NP(IBM, NNP) VP(bought, VBD)$

Assignment Project Exam Help

$\approx P_h(VP | S, bought, VBD)$ <https://powcoder.com> Generate head constituent
 $\times P_l(NP(IBM, NNP) | S, bought, VBD, VP)$ Add WeChat powcoder
 $\times P_l(NP(\text{week}, NN) | S, bought, VBD, VP)$ Generate left modifiers
 $\times P_l(STOP | S, bought, VBD, VP)$
 $\times P_r(STOP | S, bought, VBD, VP)$ Generate right modifiers

Adding other dependencies

- (Badly-named) “distance measure” to capture properties of attachment relevant to current modifier.
- $P_l(L_i(l_i) \mid X, h, H)$ becomes $P_l(L_i(l_i) \mid X, h, H, distance_i(i, 1))$ and analogously on the right.
- The value of $distance_x$ is actually a pair of Boolean random variables:
 - Is string $1..(i-1)$ of length 0?
i.e., is attachment of modifier i to the head?
 - Does string $1..(i-1)$ contain a verb?
i.e., is attachment of modifier i crossing a verb?

Collins's “model 1” 4

- Backs off ...
 - to tag probability when no data for specific word;
 - to complete non-lexicalization when necessary.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Collins's Models 2 and 3

- *Model 2:* Add verb subcategorization and argument/adjunct distinction.
- *Model 3:* Integrate gaps and trace identification into model.
Assignment Project Exam Help
- Especially important with addition of subcategorization.
<https://powcoder.com>
Add WeChat powcoder

Results and conclusions 1

- Model 2 outperforms Model 1.
- Model 3: Similar performance, but identifies traces too.
- Model 2 performs best overall:
 - LP = 89.6, LR = 89.9 [sentences \leq 100 words].
 - LP = 90.1, LR = 90.4 [sentences \leq 40 words].
- Rich information improves parsing performance.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Results and conclusions 2

- **Strengths:**

- Incorporation of lexical and other linguistic information.
- Competitive results.

- **Weaknesses:**

- Supervised training.
- Performance tightly linked to particular type of corpus used.

Results and conclusions 3

- **Importance to CL:**
 - High-performance parser showing benefits of lexicalization and linguistic information.
 - Publicly available, widely used in research.
 - There was some initial hope that it would make language models better, but that didn't pan out.
 - But it was fairly successful at giving us some access to semantics, i.e. language modelling makes parsing better.